

# Analyzing the Gender Wage Gap Over Time

Brian Mann

```
In [164]: # imported libraries
import numpy as np
import pandas as pd
import random
import matplotlib.pyplot as plt
import statsmodels.formula.api as smf
from scipy.stats import skew
from statsmodels.distributions.empirical_distribution import ECDF
from scipy.stats import lognorm, kstest
from scipy.stats import kendalltau
```

```
In [3]: # read the data into a data frame
wages = pd.read_csv("wages_by_education.csv")
wages.head()
```

```
Out[3]:
```


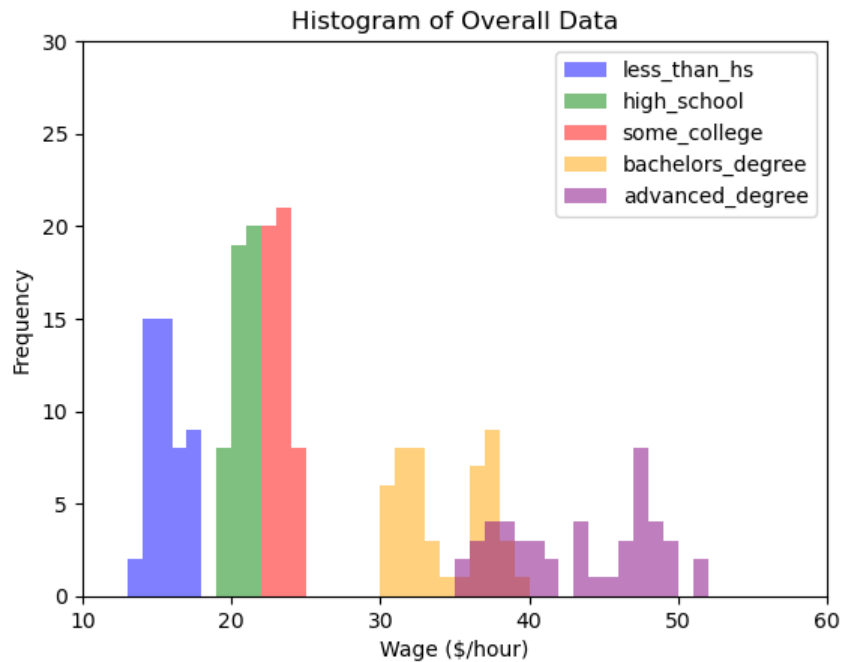
	year	less_than_hs	high_school	some_college	bachelors_degree	advanced_degree	men_less_than_hs	men_high_school	men_son
0	1973	18.06	22.22	24.08	32.80	38.16	21.18	26.90	
1	1974	17.68	21.60	23.32	31.69	38.37	20.63	26.15	
2	1975	17.30	21.55	23.30	31.45	38.41	20.00	26.02	
3	1976	17.52	21.76	23.49	31.46	37.50	20.36	26.14	
4	1977	17.59	21.50	22.97	31.07	37.36	20.43	25.97	

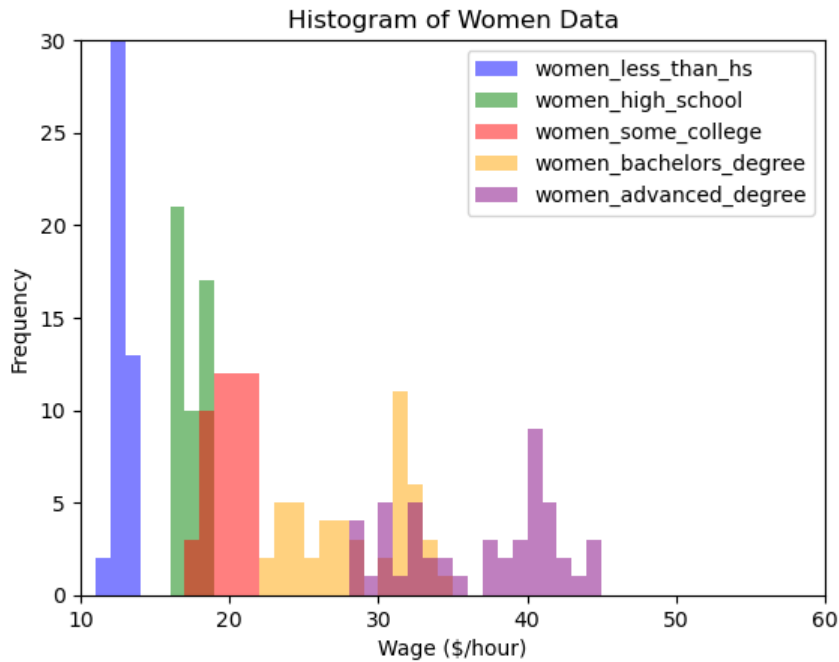
5 rows × 61 columns

```
In [10]: # variables used, separated by year, overall, men and women
years = wages.iloc[:,0]
overall = wages.iloc[:,1:6]
men = wages.iloc[:,6:11]
women = wages.iloc[:,11:16]
```

```
In [26]: # this function generates a histogram for each of the columns in the given data frame
def make_hist(df, name):
    # set up the subplots and colors
    fig, ax = plt.subplots()
    colors = ['blue', 'green', 'red', 'orange', 'purple']
    # the the binwidth to $1 and make the histogram for each column
    for i, col in enumerate(df.columns):
        ax.hist(df[col], bins=range(int(min(df[col])), int(max(df[col])) + 1, 1),
                color=colors[i], alpha=0.5, label=col)
    # make the axes limits the same for each plot
    plt.xlim(10, 60)
    plt.ylim(0, 30)
    plt.xlabel('Wage (\$/hour)')
    plt.ylabel('Frequency')
    plt.title(f'Histogram of {name} Data')
    plt.legend()
    plt.show()
```

```
In [27]: # use this list of dfs and names to generate each of the histogram plots
df_and_names = [(overall, "Overall"), (men, "Men"), (women, "Women")]
for x in df_and_names:
    df, name = x
    make_hist(df, name)
```





```
In [32]: # gather a list of each of the data frames
dfs = [overall, men, women]
```

```
In [38]: # this function takes in a data frame and gives the mean, median, range, and skew for each column
def make_summary_stats(df):
    means = df.mean()
    medians = df.median()
    ranges = df.max() - df.min()
    skews = df.apply(skew)
    summary_stats = pd.DataFrame({
        'Mean': means,
        'Median': medians,
        'Range': ranges,
        'Skew': skews
    })

    print(summary_stats)
```

```
In [40]: for df in dfs:
         make_summary_stats(df)
```

	Mean	Median	Range	Skew
less_than_hs	15.7026	15.340	4.11	0.520736
high_school	20.8766	20.855	3.08	0.134490
some_college	23.2192	23.185	3.40	0.646484
bachelors_degree	34.7686	34.205	11.61	0.319095
advanced_degree	43.8990	44.085	18.42	0.018634

	Mean	Median	Range	Skew
men_less_than_hs	17.5652	16.905	5.79	0.674019
men_high_school	23.8326	23.695	4.79	0.697004
men_some_college	26.3338	26.365	3.77	0.004734
men_bachelors_degree	39.9884	39.485	13.85	0.564821
men_advanced_degree	49.4302	48.940	24.80	0.164481

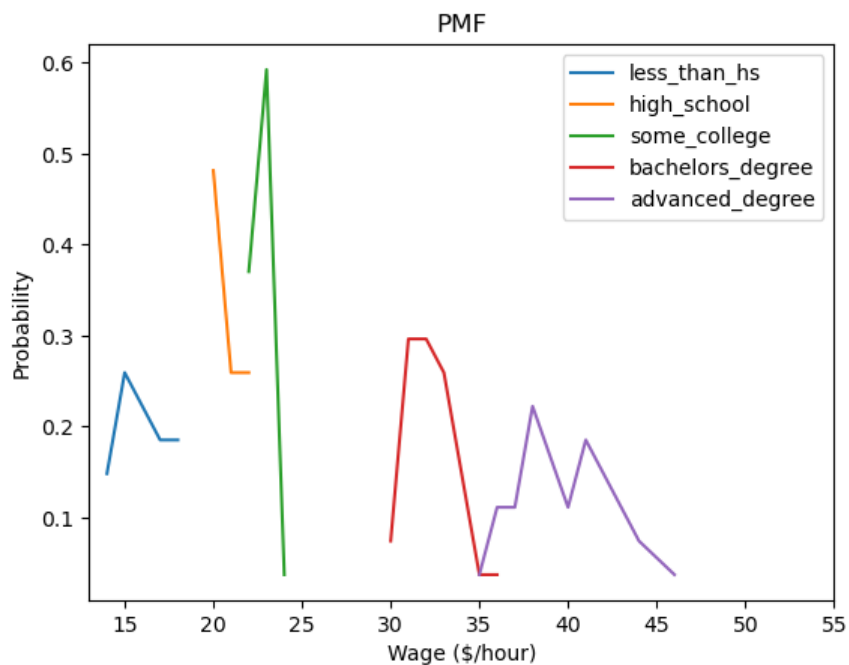
  

	Mean	Median	Range	Skew
women_less_than_hs	12.8514	12.810	2.65	0.932343
women_high_school	17.5716	17.370	2.91	0.348900
women_some_college	19.9432	19.985	4.44	-0.121558
women_bachelors_degree	28.9264	29.425	12.47	-0.145332
women_advanced_degree	36.9752	38.345	17.57	-0.206381


```
In [74]: # separate the overall data into pre-2000 and post-2000
before_2k = overall[:27].round()
after_2k = overall[27:].round()
```

```
In [109]: # this function takes in a data frame and plots the PMF for each column onto the same plot
def make_pmfs(df):
    # initialize an empty dictionary
    pmfs = dict()
    # iterate over the columns, generating a pmf via value_counts
    for col in df.columns:
        value_counts = df[col].value_counts()
        pmf = value_counts / len(df[col])
        # sort the pmf so that it is ordered
        pmf = pmf.sort_index()
        pmfs[col] = pmf
    # plot the pmf for each column
    for col, pmf in pmfs.items():
        plt.plot(pmf.index, pmf.values, label=col)
    # set consistent axes limits and labels
    plt.xlim(13, 55)
    plt.xlabel('Wage (\$/hour)')
    plt.ylabel('Probability')
    plt.title('PMF')
    plt.legend()
    plt.show()
```

```
In [110]: make_pmfs(before_2k)
```



```
In [111]: make_pmfs(after_2k)
```




```
In [88]: # make it such that every column has the same name for both the men and women data frames
men.columns = [col[4:] for col in men.columns]
women.columns = [col[6:] for col in women.columns]
```

```
In [92]: # simply subtract the two data frames to get the wage gap for each year and education level
gap = men - women
```

```
In [107]: # this function generates a line chart for each of the columns in the given data frame
def make_line_chart(df):
    for col in df.columns:
        plt.plot(years, df[col], label=col)

    plt.xlabel('Year')
    plt.ylabel('Wage Gap (\$/hour)')
    plt.title('Wage Gap Over Time')
    plt.legend()
    plt.show()
```

```
In [108]: make_line_chart(gap)
```

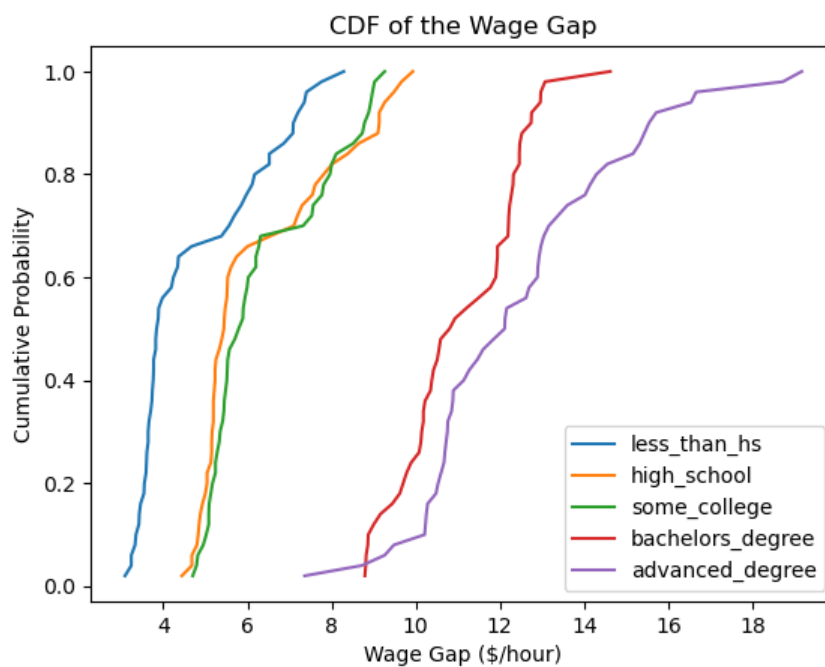


```
In [105]: # this function generates the cdf for each of the columns in the data frame and plots them together.
def make_cdf(df):
    fig, ax = plt.subplots()

    for col in df.columns:
        ecdf = ECDF(df[col])
        ax.plot(ecdf.x, ecdf.y, label=col)

    ax.set_xlabel('Wage Gap (\$/hour)')
    ax.set_ylabel('Cumulative Probability')
    ax.set_title('CDF of the Wage Gap')
    ax.legend()
    plt.show()
```

```
In [106]: make_cdf(gap)
```



```

In [165]: # this function takes in a data frame, flattens it into a single array, generates a sample log-norm
# distribution, and then plots the distribution over the given data.
def make_log_norm(df):
    # flatten all of the wage gap data into one array, then get the parameters for the
    # log-normal distribution
    wage_data = df.values.flatten()
    shape, loc, scale = lognorm.fit(wage_data)
    # get the range of the distribution
    x = np.linspace(min(wage_data), max(wage_data), 1000)
    # get the pdf of the lognormal distribution fitted to the data set
    pdf = lognorm.pdf(x, shape, loc=loc, scale=scale)
    # plot the histogram of the wage gap data, with the log-normal distribution overlayed
    plt.hist(wage_data, bins=30, density=True, alpha=0.2, color='blue', label='Wage Gap Data')
    plt.plot(x, pdf, 'r-', lw=1, label='Log-Normal Distribution')
    plt.xlabel('Wage Gap (\$/hour)')
    plt.ylabel('Probability Density')
    plt.title('Log-Normal Distribution Fit to Wage Gap Data')
    plt.legend()
    plt.show()


    ks_statistic, p_value = kstest(wage_data, 'lognorm', args=(shape, loc, scale))
    print(f"Kolmogorov-Smirnov test statistic: {ks_statistic}")
    print(f"P-value: {p_value}")

```

```

In [166]: make_log_norm(gap)

```



```


Kolmogorov-Smirnov test statistic: 0.10051534067012036
P-value: 0.011883550359003724

```


```
In [129]: # this function generates a scatter plot of the values in one dataframe vs the other data
# frame given that both have the same column names. It also generates the correlation and
# covariance
def make_scatter_plots(df1, df2):
    # generate a scatter plot over each column
    for col in df1.columns:
        plt.scatter(df1[col], df2[col], alpha=0.5)
        plt.title(f'Scatter Plot of {col} (Men vs Women)')
        plt.xlabel('Men')
        plt.ylabel('Women')
        plt.grid(True)
        plt.show()
    # get the correlation coefficient
    correlation_coefficient = np.corrcoef(df1[col], df2[col])[0, 1]
    print(f"Pearson's correlation coefficient for {col}: {correlation_coefficient}")
    # get the covariance
    covariance = np.cov(df1[col], df2[col])[0, 1]
    print(f"Covariance for {col}: {covariance}\n")
```




```
In [130]: make_scatter_plots(men, women)
```




Pearson's correlation coefficient for less\_than\_hs: 0.40936709147072853  
Covariance for less\_than\_hs: 0.3794027755102042



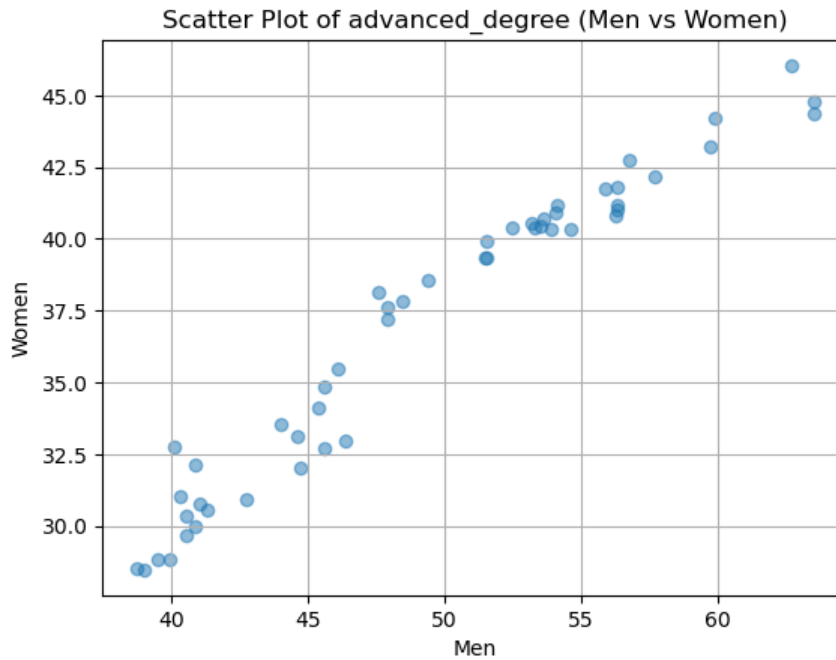
Pearson's correlation coefficient for high\_school: -0.22108317658680293  
Covariance for high\_school: -0.23151036734693842



Pearson's correlation coefficient for some\_college: 0.17219247888114797  
Covariance for some\_college: 0.19758555102040834



Pearson's correlation coefficient for bachelors\_degree: 0.9268341535241581  
Covariance for bachelors\_degree: 13.017736979591835



Pearson's correlation coefficient for advanced\_degree: 0.9740439154323426  
 Covariance for advanced\_degree: 36.14953771428571

```
In [145]: # test the hypothesis that the df values have significant changes by
# using Kendall's Tau-B
def test_hypothesis(df):
    for col in df.columns:
        tau, p_value = kendalltau(years, df[col])

        print(f"*** Hypothesis Test for {col} ***")
        print("----")
        print(f"Kendall's tau-b: {tau}")
        print(f"P-value: {p_value}")
        print("----")
```

```
In [149]: test_hypothesis(gap)

*** Hypothesis Test for less_than_hs ***
----
Kendall's tau-b: -0.5629086845967293
P-value: 8.224558841832918e-09
----
*** Hypothesis Test for high_school ***
----
Kendall's tau-b: -0.7247458917172676
P-value: 1.2388367167678222e-13
----
*** Hypothesis Test for some_college ***
----
Kendall's tau-b: -0.5482027973358568
P-value: 1.9860464050071154e-08
----
*** Hypothesis Test for bachelors_degree ***
----
Kendall's tau-b: 0.017966518443500928
P-value: 0.8539873239818369
----
*** Hypothesis Test for advanced_degree ***
----
Kendall's tau-b: 0.7142857142857143
P-value: 2.493596474326011e-13
----
```

```
In [162]: # this function takes in a data frame and runs a linear regression function with time as the predictor
# and each column as the outcome variable. It then fits the regression line to the scatter plot
# of each wage value by year.
def make_linear_regression(df):
    # iterate over each column, giving the regression data, line and scatter plot
    for col in df.columns:
        # get the linear model and generate the results of the fit
        model = smf.ols(formula=f'{col} ~ years', data=df)
        results = model.fit()
        # print a summary of the results
        print(results.summary())
        # now plot the given scatter plot and regression line
        plt.scatter(years, df[col], alpha=0.5, color="green", label='Wage Data')
        plt.plot(years, results.predict(df), color='red', label='Regression Line')
        plt.xlabel('Year')
        plt.ylabel('Wage Gap (\$/hour)')
        plt.title(f'Linear Regression for {col}')
        plt.legend()
        plt.grid(True)
        plt.show()
```


```
In [163]: make_linear_regression(gap)
```

OLS Regression Results						
Dep. Variable:	less_than_hs		R-squared:	0.686		
Model:	OLS		Adj. R-squared:	0.680		
Method:	Least Squares		F-statistic:	105.1		
Date:	Fri, 01 Mar 2024		Prob (F-statistic):	1.11e-13		
Time:	11:58:58		Log-Likelihood:	-61.668		
No. Observations:	50		AIC:	127.3		
Df Residuals:	48		BIC:	131.2		
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	174.8509	16.596	10.536	0.000	141.483	208.219
years	-0.0852	0.008	-10.252	0.000	-0.102	-0.068
Omnibus:	3.267		Durbin-Watson:	0.133		
Prob(Omnibus):	0.195		Jarque-Bera (JB):	1.646		
Skew:	-0.048		Prob(JB):	0.439		
Kurtosis:	2.116		Cond. No.	2.77e+05		

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.77e+05. This might indicate that there are strong multicollinearity or other numerical problems.



# OLS Regression Results

Dep. Variable:	high_school	R-squared:	0.748
Model:	OLS	Adj. R-squared:	0.742
Method:	Least Squares	F-statistic:	142.1
Date:	Fri, 01 Mar 2024	Prob (F-statistic):	5.92e-16
Time:	11:58:58	Log-Likelihood:	-60.818
No. Observations:	50	AIC:	125.6
Df Residuals:	48	BIC:	129.5
Df Model:	1		
Covariance Type:	nonrobust		


	coef	std err	t	P> t	[0.025	0.975]
Intercept	200.7614	16.316	12.305	0.000	167.956	233.566
years	-0.0974	0.008	-11.921	0.000	-0.114	-0.081

Omnibus:	2.466	Durbin-Watson:	0.089
Prob(Omnibus):	0.291	Jarque-Bera (JB):	1.431
Skew:	0.062	Prob(JB):	0.489
Kurtosis:	2.181	Cond. No.	2.77e+05

## Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.77e+05. This might indicate that there are strong multicollinearity or other numerical problems.



# OLS Regression Results

Dep. Variable:	some_college	R-squared:	0.657
Model:	OLS	Adj. R-squared:	0.650
Method:	Least Squares	F-statistic:	92.04
Date:	Fri, 01 Mar 2024	Prob (F-statistic):	9.67e-13
Time:	11:58:58	Log-Likelihood:	-61.548
No. Observations:	50	AIC:	127.1
Df Residuals:	48	BIC:	130.9
Df Model:	1		
Covariance Type:	nonrobust		


	coef	std err	t	P> t	[0.025	0.975]
Intercept	165.2170	16.556	9.979	0.000	131.930	198.504
years	-0.0795	0.008	-9.594	0.000	-0.096	-0.063

Omnibus:	0.573	Durbin-Watson:	0.141
Prob(Omnibus):	0.751	Jarque-Bera (JB):	0.659
Skew:	0.005	Prob(JB):	0.719
Kurtosis:	2.438	Cond. No.	2.77e+05

## Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.77e+05. This might indicate that there are strong multicollinearity or other numerical problems.



# OLS Regression Results

Dep. Variable:	bachelors_degree	R-squared:	0.021
Model:	OLS	Adj. R-squared:	0.001
Method:	Least Squares	F-statistic:	1.029
Date:	Fri, 01 Mar 2024	Prob (F-statistic):	0.315
Time:	11:58:58	Log-Likelihood:	-87.930
No. Observations:	50	AIC:	179.9
Df Residuals:	48	BIC:	183.7
Df Model:	1		
Covariance Type:	nonrobust		


	coef	std err	t	P> t	[0.025	0.975]
Intercept	39.5306	28.061	1.409	0.165	-16.889	95.950
years	-0.0143	0.014	-1.015	0.315	-0.042	0.014

Omnibus:	0.363	Durbin-Watson:	0.263
Prob(Omnibus):	0.834	Jarque-Bera (JB):	0.528
Skew:	0.149	Prob(JB):	0.768
Kurtosis:	2.595	Cond. No.	2.77e+05

## Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.77e+05. This might indicate that there are strong multicollinearity or other numerical problems.





### OLS Regression Results

Dep. Variable:	advanced_degree	R-squared:	0.746
Model:	OLS	Adj. R-squared:	0.741
Method:	Least Squares	F-statistic:	141.1
Date:	Fri, 01 Mar 2024	Prob (F-statistic):	6.70e-16
Time:	11:58:58	Log-Likelihood:	-81.286
No. Observations:	50	AIC:	166.6
Df Residuals:	48	BIC:	170.4
Df Model:	1		
Covariance Type:	nonrobust		


	coef	std err	t	P> t	[0.025	0.975]
Intercept	-279.4313	24.569	-11.373	0.000	-328.831	-230.032
years	0.1461	0.012	11.881	0.000	0.121	0.171

Omnibus:	1.731	Durbin-Watson:	0.772
Prob(Omnibus):	0.421	Jarque-Bera (JB):	1.510
Skew:	0.418	Prob(JB):	0.470
Kurtosis:	2.838	Cond. No.	2.77e+05

#### Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.  
[2] The condition number is large, 2.77e+05. This might indicate that there are strong multicollinearity or other numerical problems.



In [ ]: